



DEPARTMENT OF COMPUTER SCIENCE,
FACULTY OF MATHEMATICS, PHYSICS AND INFORMATICS,
COMENIUS UNIVERSITY IN BRATISLAVA

DISTANCE ORACLES FOR TIMETABLE GRAPHS

(Master thesis)

bc. František Hajnovič

Study program: Computer science

Branch of study: 2508 Informatics

Supervisor: doc. RNDr. Rastislav Kráľovič, PhD.

Bratislava 2013



Comenius University in Bratislava
Faculty of Mathematics, Physics and Informatics

THESIS ASSIGNMENT

Name and Surname: Bc. František Hajnovič
Study programme: Computer Science (Single degree study, master II. deg., full time form)
Field of Study: 9.2.1. Computer Science, Informatics
Type of Thesis: Diploma Thesis
Language of Thesis: English
Secondary language: Slovak

Title: Distance oracles for timetable graphs

Aim: The aim of the thesis is to explore the applicability of results about distance oracles to timetable graphs. It is known that for general graphs no efficient distance oracles exist, however, they can be constructed for many classes of graphs. Graphs defined by timetables of regular transport carriers form a specific class which it is not known to admit efficient distance oracles. The thesis should investigate to which extent the known desirable properties (e.g. small highway dimension) are present in these graphs, and/or identify new ones. Analytical study of graph operations and/or experimental verification on real data form two possible approaches to the topic.

Supervisor: doc. RNDr. Rastislav Kráľovič, PhD.
Department: FMFI.KI - Department of Computer Science
Vedúci katedry: doc. RNDr. Daniel Olejár, PhD.
Assigned: 08.11.2011

Approved: 15.11.2011
prof. RNDr. Branislav Rován, PhD.
Guarantor of Study Programme

Student

Supervisor



Univerzita Komenského v Bratislave
Fakulta matematiky, fyziky a informatiky

ZADANIE ZÁVEREČNEJ PRÁCE

Meno a priezvisko študenta: Bc. František Hajnovič
Študijný program: informatika (Jednoodborové štúdium, magisterský II. st., denná forma)
Študijný odbor: 9.2.1. informatika
Typ záverečnej práce: diplomová
Jazyk záverečnej práce: anglický
Sekundárny jazyk: slovenský

Názov: Efektívny výpočet vzdialeností v grafoch spojení lineík.

Cieľ: Cieľom práce je preštudovať možnosti aplikácie výsledkov o distance oracles v grafoch reprezentujúcich dopravné siete na grafy spojení lineík. Otázka, či a aké dôležité vlastnosti ostávajú zachované sa dá riešiť teoreticky pre rôzne triedy grafov a/alebo experimentálne pre reálne dáta.

Vedúci: doc. RNDr. Rastislav Kráľovič, PhD.

Katedra: FMFI.KI - Katedra informatiky

Vedúci katedry: doc. RNDr. Daniel Olejár, PhD.

Dátum zadania: 08.11.2011

Dátum schválenia: 15.11.2011

prof. RNDr. Branislav Rován, PhD.
garant študijného programu

študent

vedúci práce

I hereby declare that I wrote this thesis by myself, only with the help of the referenced literature,
under the careful supervision of my thesis advisor.

.....

Acknowledgements

I would like to thank ...

Abstract

This thesis...

Key words: **oracles**, **timetable**

Abstrakt

V této práci...

Klíčové slova: **oracles**, **timetable**

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Approach	1
1.3	Goals	1
1.4	Organization	1
2	Preliminaries	2
2.1	Objects	2
2.2	Earliest arrival and optimal connection	5
2.3	(Distance) Oracles	6
2.4	Dijkstra’s algorithm	7
3	Related work	8
4	Data & analysis	9
4.1	Data	9
4.2	Basic properties	10
5	Underlying shortest paths	12
5.1	<i>USP-OR</i>	13
5.1.1	Analysis of <i>USP-OR</i>	15
5.1.2	Performance of <i>USP-OR</i>	17
5.2	<i>USP-OR-A</i>	17
5.2.1	Analysis of <i>USP-OR-A</i>	20
5.2.2	Correctness of <i>USP-OR-A</i>	23
5.2.3	Modifications of <i>USP-OR-A</i>	24
5.3	Selection of access node set	24
5.3.1	Choosing the optimal access node set	25
5.3.2	Choosing ANs based on node properties	27
5.3.3	Choosing ANs heuristically	28
5.4	Performance of <i>USP-OR-A</i> and comparisons	28
5.4.1	<i>USP-OR-A</i> with <i>Locsep</i>	28
5.4.2	<i>USP-OR-A</i> with <i>Locsep Max</i>	28
6	Neural network approach	29
7	Application TTBlazer	30
8	Conclusion	31
	Appendix A File formats	32

1 Introduction

World is getting smaller every day...

1.1 Motivation

1.2 Approach

1.3 Goals

1.4 Organization

2 Preliminaries

In this section, we provide most of the definitions and terminology used throughout the thesis.

2.1 Objects

First, we will formalize the notion of a timetable and its derived graph forms, the underlying graph and terms related to these objects.

Definition 2.1. Timetable (TT)

A timetable is a set $T = \{(x, y, p, q) \mid p, q \in \mathbb{N}, p < q\}$.

- Elements of T (the 4-tuples) are called **elementary connections**. For an elementary connection $e = (x, y, p, q)$:
 - $from(e) = x$ is the **departure city**
 - $to(e) = y$ is the **arrival/destination city**
 - $dep(e) = p$ is the **departure time**
 - $arr(e) = q$ is the **arrival time**
- The set of all **cities** will be denoted as $ct_T = \{x \mid (x, y, p, q) \in T \text{ or } (y, x, p, q) \in T\}$ and the number of cities as n_T
- Pairs (x, p) or (y, q) such that $(x, y, p, q) \in T$ form the set of **events** ev_T . The set of events in a specific city x is $ev_T(x) = \{(x, t) \mid (x, y, t, q) \in T \text{ or } (y, x, p, t) \in T\}$
- Let $tlow_T = \min_{e \in T} dep(e)$ and $thigh_T = \max_{e \in T} arr(e)$. The value $r_T = thigh_T - tlow_T$ is called the **time range** of the timetable.
- **Height** of the timetable is the maximum number of events in a city: $h_T = \max_{x \in cities_T} \{|ev_T(x)|\}$

Let us describe some the defined terms more informally. An elementary connection corresponds to moving from one stop to the next one, e.g. with a bus (thus we disregard the notion of *lines*, i.e. getting on and off). Note that we express time as an integer - throughout this paper, this integer will represent the minutes elapsed from the time 00:00 of the first day. Thus we may take the liberty of talking about time in integer or *days hh:mm* format, as convenient at the moment. Lastly, an event simply represent an arrival or departure of a e.g. train at some station. The remaining terms should be clear enough.

Place		Time	
From	To	Departure	Arrival
A	B	10:00	10:45
B	C	11:00	11:30
B	C	11:30	12:10
B	A	11:20	12:30
C	A	11:45	12:15

Table 2.1: An example of a timetable - the set of elementary connections (between pairs of **cities**). An example of an event is a pair (A, 10:00), when some el. connection departs from A.

Following is a definition of a connection.

Definition 2.2. Connection

A connection from a to b is a sequence of elementary connections $c = (e_1, e_2, \dots, e_k), k \geq 1$, such that $from(e_1) = a$, $to(e_k) = b$ and $\forall i \in \{2, \dots, k\} : (to(e_i) = from(e_{i-1}), arr(e_i) \geq dep(e_{i-1}))$.

- Connection **starts** at the departure time $\text{start}(\mathbf{c}) = \text{dep}(e_1)$ and **ends** at the arrival time $\text{end}(\mathbf{c}) = \text{arr}(e_k)$.
- We also extend $\text{from}(\mathbf{c}) = \text{from}(e_1)$ and $\text{to}(\mathbf{c}) = \text{to}(e_k)$
- **Length** of the connection is $\text{len}(\mathbf{c}) = \text{end}(\mathbf{c}) - \text{start}(\mathbf{c})$
- **Size** of the connection is $\text{size}(\mathbf{c}) = k$ ¹
- We will denote the set of **all connections** from a to b in a timetable T as $\mathbf{C}_T(a, b)$. We also define $\mathbf{C}_T = \cup_{a,b} \mathbf{C}_T(a, b)$

So we understand connection as a (valid) sequence of elementary connections.




Figure 2.1: A valid connection made out of **elementary connections** (and **waiting**, which is implicit).

Next, we continue with the underlying graph - a graph representing basically the map on top of which the timetable operates.

Definition 2.3. Underlying graph (UG graph)

The underlying graph of a timetable T , denoted \mathbf{ug}_T , is an oriented graph (V, E) , where V is the set of all timetable cities and $E = \{(x, y) \mid \exists (x, y, p, q) \in T\}$

- By \mathbf{m}_T we will denote the number of arcs in the UG

Note, that we do not specify the weights of the edges in the underlying graph - they will be specified based on the current usage of the UG. Most of the time, however, if we work with a weighted UG, the weight of an arc will be the length of the shortest elementary connection on that arc. More specifically, $w(x, y) = \min_{(x, y, p, q) \in T} (q - p) \ \forall (x, y) \in E(\mathbf{ug}_T)$. Such weighted UG will be called **optimistic** (denoted $\mathbf{ug}_T^{\text{opt}}$).




Figure 2.2: An optimistic underlying graph of the timetable 2.1. The nodes are the **cities** of the timetable.

¹We will use similar terminology when talking about paths - the *size* is the number of vertices (hops) in the path while the *length* refers to the actual distance (sum of weights of the edges in the path)

If we want to represent the timetable by a graph, there are two most common options [?] - the time-expanded and time-dependent graph.

Definition 2.4. Time-expanded graph (TE graph)

Let T be a timetable. Time-expanded graph from T , denoted te_T , is an oriented graph (V, E) whose vertices correspond to events of T , that is $V = \{(x, t) \mid (x, t) \in ev_T\}$. The edges of G are of two types

1. $([x, p], [y, q]) \forall (x, y, p, q) \in T$ - the so called **connection edges**
2. $([x, p], [x, q]) \mid [x, p], [x, q] \in V, p < q$ and $\nexists [x, r] \in V : p < r < q$. - the so called **waiting edges**

Weight of the edge $([x, p], [y, q])$ is $w([x, p], [y, q]) = q - p$.

Informally, an edge in TE graph represent either the travelling with an elementary connection or waiting for the next event in the same city. Also, the time range and height of a timetable could be easily illustrated on the TE graph (see picture 2.3).




Figure 2.3: Time-expanded graph of the timetable 2.1. Nodes represent the **events**. There are **connection** and **waiting** edges (dashed). The time range is 2h:30m and the height is 4 (as there are 4 events in city B).

Definition 2.5. Time-dependent graph (TD graph)

Let T be a timetable. Time-dependent graph from T , denoted td_T , is an oriented graph (V, E) whose vertices are the timetable cities and $E = \{(x, y) \mid \exists (x, y, p, q) \in T\}$. Furthermore, the weight of an edge $(x, y) \in E$ is a piece-wise linear function $w(x, y) = f_{x,y}(t) = q - t$ where q is:

- $\min\{arr(e) \mid e \in T, dep(e) \geq t\}$
- ∞ , if $dep(e) < t \forall e \in T$

Intuitively, the TD graph is simply the UG graph where each arc carries a function specifying the traversal time of that arc at any time. For an example, see picture 2.5: The latest point of every linear segment is called the **interpolation point** and it corresponds to an elementary connection (its coordinates are $dep(e), len(e)$ for corresponding el. connection e). Note that a list of all interpolation points fully defines the piece-wise linear function.

The algorithms in this thesis use almost exclusively the TD graphs, mainly because they are less space consuming. Also, time-dependent Dijkstra searches are a bit faster on TD graphs,

because the search space that has to be explored is smaller. On the other hand, TE graphs are more flexible when we need to take additional search parameters into consideration (like transfers, travel costs). Since we will not talk about these, TD graphs are more suitable.




Figure 2.4: Time-dependent graph of the timetable 2.1. The nodes are the cities.




Figure 2.5: Piece-wise linear function - traversal times for the arc (B, C) . The highlighted points are the interpolation points.

To sum up, there are four main types of objects we will be working with:

- Timetable (TT)
- Underlying graph (UG)
- Time-expanded graph (TE)
- Time-dependent graph (TD)

For further reference, we will call **timetable objects** those, that fully represent a timetable (TT, TE, TD) and **graph objects** those, that can be viewed as a graph (UG, TE, TD).

Note: Throughout this paper, we will relax a bit the notation and leave out subscripts (e.g. $ug_T \rightarrow ug$, $n_T \rightarrow n$, etc.) in situations, where the context is clear enough.

2.2 Earliest arrival and optimal connection

Now we would like to formulate the main problems this thesis deals with.

Definition 2.6. Earliest arrival problem (EAP)

Given a timetable T , departure city x , destination city y and a departure time t , the task is to determine $\mathbf{t}_{(x,t,y)}^* = \min_{c \in C_T(x,y)} \{t + \text{len}(c) | \text{start}(c) \geq t\}$.

- We will refer to the tuple (x, t, y) as an **EAP instance**, or an **EAP query**
- The time $\mathbf{t}_{(x,t,y)}^*$ is called the **earliest arrival (EA)** for the given EAP instance

A bit more difficult version of this problem is one, where we require to actually output the connection ending at time given by EA.

Definition 2.7. Optimal connection problem (OCP)

Given a timetable T , departure city x , destination city y and a departure time t , the task is to determine the **optimal connection (OC)** $\mathbf{c}_{(a,t,b)}^* = \text{argmin}_{c \in C_T(a,b)} \{t + \text{len}(c) | \text{start}(c) \geq t\}$.

The instance/query in case of the optimal connection problem has the same form as EAP query. Also, note that the OCP is at least as hard to solve as EAP since having the optimal connection implies the optimal (earliest) arrival time.. In order to avoid technical issues in later parts of the thesis, we will assume the optimal connection is unique (i.e., there is not a different connection with the same end time) or that ties are won by a lexicographically first connection.

Example 2.1. Consider our timetable from table 2.1. For the EAP instance $(B, 10:45, A)$, the earliest arrival (EA) is 12:15 and the optimal connection (OC) is $((B, C, 11:00, 11:30), (C, A, 11:45, 12:15))$, as could be easily seen from picture 2.6 of the TE graph.




Figure 2.6: Optimal connection and earliest arrival time are marked in **bold**.

2.3 (Distance) Oracles

The term *distance oracle* was first coined in 2001 by Thorup and Zwick [?], when talking about quick shortest path (or distance) computations on graphs. One approach to this problem is to pre-compute some information on the graph to speed-up answering of the queries. The paper of Thorup and Zwick was dealing with trade-offs among the time complexity of the pre-computation, the amount of pre-computed information, the speed-up in query times and the accuracy of the answers. Since the pre-computed data structure is something that helps us answer the queries more efficiently, it resembles an oracle, thus the term distance oracle.

In this thesis, we will discuss methods that behave the same way, but deal with the earliest arrival problem (or optimal connection problem) - there is some pre-processing of the timetable with a resulting data structure that speeds up answering subsequent queries. To formalize this a little more, we will refer to this kind of methods as **oracle based methods**. For such a method m , we are interested mainly in its four parameters:

- **Preprocessing time** ($prep(m)$) - the time complexity of the pre-computation
- **Preprocessed space** ($size(m)$) - the space complexity of the pre-computed data structure (the so called **oracle**)

- **Query time** ($qtime(m)$) - the time complexity of answering a single query
- **Stretch** ($stretch(m)$) - the worst-case ratio against the optimal value of earliest arrival (the lower, the better)

The preprocessing time is probably the least critical resource. A reasonable polynomial should bind its time complexity, depending on the computational power of the user and the scale of the timetable. The size of the preprocessed oracle is much more important - in the optimal case, it should be bound by the space complexity of the timetable itself. Optimality of the query time depends on which problem we are solving. If we query for the whole optimal connection, we have to count with a time complexity at least proportional to the diameter of the underlying graph (as connections could be that long, or even longer). If we require only the EA value as an output, much better speed-ups could be expected. The stretch should be of course as low as possible.

2.4 Dijkstra's algorithm

Throughout this thesis, we will often use Dijkstra's algorithm and its modifications both as a part of our algorithms and as a reference point against which we will compare the performance of our methods. This is a common practice. Researchers working on methods answering distance or shortest path queries in road networks commonly use the term *speed-up*, i.e. *how many times faster* is their algorithm against the Dijkstra's algorithm.

Dijkstra's algorithm is originally an algorithm that looks for shortest paths in weighted oriented graphs. It was published by E. W. Dijkstra in 1959 [?] and we will not explain it at this place, as the algorithm is very well explained at many other places (e.g. [?]). For a good summary of Dijkstra's algorithm related implementations and publications see [?].

As our task is to compute earliest arrivals or optimal connections instead of distances and shortest paths, our "reference point" will be a slightly modified Dijkstra's algorithm called **time-dependent Dijkstra's algorithm** ?? (or TD Dijkstra for short). The algorithm is run on a time-dependent graph and works just like the ordinary Dijkstra's algorithm, except that the weight of each arc (x, y) is determined for the time t at which we had settled vertex x .

If we assume that the evaluation of an arc by the cost function of the TD graph is implemented in constant time, the running time of the TD Dijkstra is $\mathcal{O}(n^2)$, just like the normal Dijkstra's algorithm. On sparse graphs, this bound can be improved using a quick data structure to determine the next node we settle. A good option is a priority queue implemented as a *Fibonacci heap*, which implements deletion in $\mathcal{O}(\log n)$ and all other operations in constant amortized time [?]. This yields the running time of TD Dijkstra $\mathcal{O}(n \log n + m)$.

We may therefore introduce a fifth parameter of our oracle based methods, the speed-up:

Definition 2.8. *Speed-up* ($spd(m)$)

A *speed-up* of an oracle based method m is the ratio $\frac{qtime_{avg}(TDDijkstra)}{qtime_{avg}(m)}$ where $qtime_{avg}(m')$ is the average query time of the respective oracle based method m' ².

The definition is rather loose in the sense that we may refer to a concrete speed-up of the method on a concrete dataset or a general, theoretical speed-up expressed as a function of the size of input.

²Note that we may also consider the TD Dijkstra algorithm to be an oracle based method - it just happens that it does not require any preprocessing.

3 Related work

4 Data & analysis

In this section we would like to introduce the timetable datasets we were working with and provide the results of the analysis which we carried out on the data. The main reason for this analysis is that it gives some insight into the properties of the timetables, and thus may contribute to the make an oracle based method with better qualities.

4.1 Data

We have obtained timetable datasets from numerous sources, in varying formats and of different types. Some of them were freely available on the Internet while others were provided by companies upon demand. Let us briefly describe each of these timetables.

The dataset *air01* contains schedules of **domestic flights in United States** for the January of 2008. It is not comprehensive in the sense that it contains entries only for flights of some of the major airports in US. However it is large enough for our purposes (almost 300 airports). This dataset is just a fraction of the data that are freely available at the pages of American Statistical Association ³ in CSV format.

Timetables *cpzu* and *cpza* represent the **regional bus** schedules from the areas of **Ružomberok and Žilina, Slovakia**. The data were provided by the company in charge of the *cp.sk* portal - In-prop s.r.o. . Both of the timetables concern about 1000 bus stops and came in a JDF 1.9 format ⁴. Apart from the actual schedules, the data in JDF contain numerous other information, which were not relevant for our purposes. From both timetables, we have extracted subsets with a time range of one day.

The *montr* dataset is part of a public feed for **Greater Montreal public transportation**, available at Google Transit Feeds ⁵. The data are in a GTFS format (defines relations between CSV files listing stations, routes, stop-times...) and were made available by Montreal's Agence métropolitaine de transport. Our timetable *montr* corresponds to daily schedules of the Chambly-Richelieu-Carignan bus services (more than 200 bus stops).

Also in GTFS format come the data of **French railways** operated by company SNCF, publicly available at their website ⁶. The schedules are weekly, but we have extracted just a sub-range corresponding to Monday. Also, there were two types of schedules: one for intercity trains and one for TER trains (regional trains). Thus the three timetables *snCF-inter* (366 stations), *snCF-ter* (2637 stations) and their union *snCF* (2646 stations).

Finally, one more country-wide railway timetable was provided by ŽSR, the company in charge of the **Slovak national railways**. This timetable was exported in a MERITS format and its time range is for one year. The number of stations in *zsr* dataset is 233.

With the help of Python and Bash scripts, we converted each of these datasets to our timetable format (described in appendix A). This timetables were then loaded by our application TTBlazer and sub-timetables (with less stations, smaller time-range or smaller height) were generated. Also the UG, TE and TD were generated from each timetable.

For a summary of the used timetables' descriptions, see table 4.1 and for their main properties, refer to table 4.2.

³<http://stat-computing.org/dataexpo/2009/the-data.html>

⁴Jednotný datový formát (JDF)

⁵<http://code.google.com/p/googletransitdatafeed/wiki/PublicFeeds>

⁶<http://test.data-sncf.com/index.php/ter.html>

Name	Description	Format	Provided by	Publicly available
<i>air01</i>	domestic flights (US)	CSV	American Stat. Assoc.	✓
<i>cpru</i>	regional bus (Ružomberok, SVK)	JDF 1.9	Inprop s.r.o.	✗
<i>cpza</i>	regional bus (SVK, Žilina)	JDF 1.9	Inprop s.r.o.	✗
<i>montr</i>	public transport (Montreal, CA)	GTFS	Montreal AMT	✓
<i>sncf</i>	country-wide intercity rails (FRA)	GTFS	SNCF	✓
<i>zsr</i>	country-wide rails (SVK)	MERITS	ŽSR	✗

Table 4.1: Timetable descriptions.

Name	El. conns.	Cities	UG arcs	Time range	Height
<i>air01</i>	601489	287	4668	1 month	24374
<i>cpru</i>	37148	871	2415	1 day	239
<i>cpza</i>	60769	1108	2778	1 day	370
<i>montr</i>	7153	217	349	1 day	363
<i>sncf</i>	90676	2646	7994	1 day	488
<i>sncf-inter</i>	4796	366	901	1 day	209
<i>sncf-ter</i>	85932	2637	7647	1 day	488
<i>zsr</i>	932052	233	588	1 year	60308

Table 4.2: Main properties of the timetables. The value of time range is approximate.

Some of the timetables have time range greater then 1 day. Furthermore, even for those marked as a 1 day timetable the exact time range is different (e.g., part of the Monday timetable might be some overnight trains with arrival on Tuesday morning). To see better the differences in the properties of different timetable types (train, flight, bus...), we made sub-timetables with 200 cities and with the upper bound on time range 1 day, 6 hours ⁷ ($thigh_T < 1 \text{ day}, 6h \forall T$) from each of our dataset. See table 4.3 for details.

Name	El. conns.	Cities	UG arcs	Exact time range	Height
<i>air01-200d</i>	19546	200	3986	1 day, 05h:00m	766
<i>cpru-200d</i>	8721	200	647	0 days, 18h:45m	239
<i>cpza-200d</i>	13225	200	583	0 days, 19h:01m	370
<i>montr-200d</i>	6985	200	320	0 days, 20h:33m	363
<i>sncf-200d</i>	8599	200	601	1 day, 05h:29m	456
<i>sncf-inter-200d</i>	2283	200	466	1 day, 01h:10m	186
<i>sncf-ter-200d</i>	7617	200	585	1 days, 00h:02m	450
<i>zsr-200d</i>	2289	200	464	1 day, 03h:26m	142

Table 4.3: 200-station sub-timetables with the maximal time range of little more than one day.

Also, to provide idea as to how big the time-expanded graphs can get consult table ??.

4.2 Basic properties

First we will take a look at the optimal connection *sizes* (size is the number of el. connections) in the timetables. For a given timetable T , we will denote the average optimal connection size as γ_T

⁷We took all elementary connections that were within our time range. From this timetable, we made an UG and its (random) sub-graph of 200 cities. Finally we selected only those elementary connections, that were on top of this sub-graph to form a timetable with 200 cities and the desired time range

and will call it the **optimal connection radius** (OC radius). We computed an approximate OC radius for each of our datasets by measuring an average connection size of sufficiently many OCs. The results in table 4.4 indicate that the average OC sizes move around value \sqrt{n} (with exception of the *air01* timetable).

Next we would like to get an idea of the sparsity of the underlying graphs. We see from the table 4.2 that the graphs are pretty sparse (again, with exception of *air01*), but we would like to make sure that the sparsity is uniform. More specifically, we will be interested in the δ -density:

Definition 4.1. δ -density

A graph G of n vertices and m arcs is δ -dense $\iff \forall G' \subseteq G, n' \geq \sqrt[4]{n} : \frac{m'}{n'} \leq \delta$

- For a timetable T , we will denote its **density** as $\delta_T = \min\{\delta \mid \text{ug}_T \text{ is } \delta\text{-dense}\}$

To find out at least approximate δ_T values for our timetables, we have randomly sampled their UGs for (connected) sub-graphs of various sizes (starting from $\sqrt[4]{n}$). In table 4.5 you can see the maximal density found during the sampling.

Name	γ_T
<i>air01</i>	2.4
<i>cpru</i>	32.3
<i>cpza</i>	33.2
<i>montr</i>	20.6
<i>sncf</i>	25.8
<i>sncf-inter</i>	8.3
<i>sncf-ter</i>	27.5
<i>zsr</i>	15.1

Table 4.4: OC radius.

Name	Maximal δ_T found
<i>air01</i>	34.5
<i>cpru</i>	4.0
<i>cpza</i>	3.4
<i>montr</i>	1.9
<i>sncf</i>	4.3
<i>sncf-inter</i>	3.1
<i>sncf-ter</i>	4.7
<i>zsr</i>	3.2

Table 4.5: Approximate density of the underlying graphs.

5 Underlying shortest paths

In section 2 we have defined a timetable as a set of elementary connections. While do not pose any other restrictions on this set or on the elementary connections themselves, the real world timetables usually have a specific nature. Quite often are the connections repetitive, that is, the same sequence of elementary connections is repeated in several different moments throughout the day.

Another thing we may notice is that if we talk about *optimal* connections between a pair of distant cities u and v , we are often left with a few possibilities as to *which way should we go*. This is not only because the underlying graph is usually quite sparse ⁸, but also because for longer distances we generally need to make use of some express connection that stops only in (small number of) bigger cities.

Thus the main idea which will repeat often throughout this section: *when carrying out an optimal connection between a pair of cities, one often goes along the same path regardless of the starting time*.

To formalize this idea, we will introduce the definition of an *underlying shortest path* - a path in UG that corresponds to some optimal connection in the timetable. To do this, we will first define a function *path* that extracts the **underlying path** (trajectory in the UG) from a given connection. Let c be a connection $c = (e_1, e_2, \dots, e_k)$.

$$\mathbf{path}(c) = \mathit{shrink}(\mathit{from}(e_1), \mathit{from}(e_2), \dots, \mathit{from}(e_k), \mathit{to}(e_k))$$

Note, that if the connection involves waiting in a city (as e.g. in picture 5.1), $e_x^i = e_x^{i+1}$ for some i . That is why we apply the *shrink* function, which replaces any sub-sequences of the type (z, z, \dots, z) by (z) in a sequence. This was rather technical way of expressing a simple intuition - for a given connection, the *path* function simply outputs a sequence of visited cities. Now we can formalize the underlying shortest path.

Definition 5.1. Underlying shortest path (USP)

A path $p = (v_1, v_2, \dots, v_k)$ in UG_T is an **underlying shortest path** if and only if $\exists t \in \mathbb{N} : p = \mathbf{path}(c_{(v_1, t, v_k)}^*), c_{(v_1, t, v_k)}^* \in C_T$




Figure 5.1: The *path* function applied on a connection to get the underlying path.

Please note that the terminology might be a bit misleading - an USP is not necessarily a shortest path in the given UG. Connections on a shortest path may simple require too much waiting (the

⁸Maybe with exception of the airline timetables, which tend to be more dense

el. connections simply do not follow well enough one another) and thus it might be that travelling along the paths with greater distance proof to be faster options.

5.1 USP-OR

We can easily extract the underlying path from a given connection. Now let us look at this from the other way - if, for a given EA query, we know the underlying shortest path, can we reconstruct the optimal connection? One thing we could do is to blindly follow the USP and at each stop take the first elementary connection to the next stop on the USP. This simple algorithm called *Expand* is described in algorithm 1.

Algorithm 1 Expand

Input

- timetable T
- path $p = (v_1, v_2, \dots, v_k)$, $v_i \in ct_T$
- departure time t

Algorithm

```

 $c$  = empty connection
 $t' = t$ 
for all  $i \in \{1, \dots, k - 1\}$  do
     $e = \operatorname{argmin}_{e' \in C_T(v_i, v_{i+1})} \{dep(e') \mid dep(e') \geq t'\}$       # take first available el. conn.
     $t' = arr(e)$ 
     $c := e$       # add the el.conn to the resulting connection
end for

```

Output

- connection c
-

Will we get an optimal connection if we expanded all possible USPs between a pair of cities? We show that we will, provided the timetable has no *overtaking* of elementary connections.

Definition 5.2. Overtaking

An elementary connection e_1 **overtakes** e_2 if, and only if $dep(e_1) > dep(e_2)$ and $arr(e_1) < arr(e_2)$.

Lemma 5.1. Let T be a timetable without overtaking, (x, t, y) an EA query in this timetable and $\mathcal{P} = \{p_1, p_2, \dots, p_k\}$ a set of all USPs from x to y . Define $c_i = \text{Expand}(T, p_i, t)$ to be the connection returned by the algorithm Expand 1. Then $\exists j : c_j = c_{x,t,y}^*$.

Proof. The optimal connection $c_{x,t,y}^*$ has an USP p which must be present in the set \mathcal{P} , as it is the set of all USPs from x to y . So $p = p_j = (v_1, v_2, \dots, v_l)$ from some j . We want to show that c_j is the optimal connection. This may be shown inductively:

1. *Base:* Expand reaches city $v_1 = x$ as soon as possible (since the connection just starts there)
2. *Induction:* Expand reached city v_i as soon as possible, it then takes the first available el. connection to the next city v_{i+1} . Since the el. connections do not overtake, Expand reached the city v_{i+1} as soon as possible.

□

We would like to stress that overtaking is understood as a situation when one carrier overtakes another between *two subsequent stations*. This situation is not that common, however it is still present in the real world timetables⁹, as shown in table 5.1. All the same, we can simply remove the

⁹In Slovak rails, no overtaking has been detected. This is not surprising as (to my knowledge) there are no inter-station tracks with multiple rails going in one direction. French railways, on the other hand have designated high-speed tracks and thus overtaking is not impossible.




Figure 5.2: An example of **overtaking** (in thick), depicted in a TE graph.

overtaken el. connections from the timetables, as they can be substituted by the quicker connection plus some waiting.

Name	Overtaken edges (%)
<i>air01</i>	1%
<i>cpru</i>	2%
<i>cpza</i>	2%
<i>montr</i>	1%
<i>sncf</i>	2%
<i>sncf-ter</i>	2%
<i>sncf-inter</i>	8%
<i>zsr</i>	0%

Table 5.1: Presence of overtaking in the timetables.

The basic idea of the algorithm *USP-OR* (a short-cut for USP oracle) is therefore simply to pre-compute all the USPs for each pair of cities. Upon a query, the algorithm simply expands all the USPs for a given pair of cities, reconstructs respective connections and chooses the best one.

Algorithm 2 *USP-OR* query

Input

- timetable T
- OC query (x, t, y)

Pre-computed

- $\forall x, y$: set of USPs between x and y ($usps(x, y)$)

Algorithm

```
 $c^* = null$ 
for all  $p \in usps_{x,y}$  do
   $c = Expand(T, p, t)$ 
   $c^* = \text{better out of } c^* \text{ and } c$ 
end for
```

Output

- connection c
-

5.1.1 Analysis of *USP-OR*

We will now have a look at the four parameters of this oracle based method. As for the preprocessing time, we need to find optimal connections from each *event* in the timetable to each *city* (or in other words - solve all possible OC queries). On these connections we apply the *path* function to obtain the USPs. The maximum number of events in one city is the height h and there is n cities, thus hn is the upper bound on the number of events. One search from a single event to all cities can be done in time $\mathcal{O}(n \log n + m)$ with a TD Dijkstra's algorithm run on the time-dependent graph of our timetable (TD_T). In worst case, m could be as much as n^2 but we may bound it as $m \leq \delta_T n$ (where δ_T is the sparsity of the timetable, defined in section 4). We therefore get the **preprocessing time** $\mathcal{O}(hn^2(\log n + \delta))$.

As for the preprocessed space, we need to store USPs for each pair of the cities (n^2 pairs) and each USP might be long at most $\mathcal{O}(n)$ hops. What is more, there might be many USPs for a single pair of cities. Therefore we have two questions with respect to the space complexity of the preprocessing:

1. What is the average size of the USPs?
2. How many are there USPs between a single pair of cities?

The answer for the first question is that the average USP size is equal to the OC radius of the timetable (γ_T) defined in section 4. An upper bound for this value is $\mathcal{O}(n)$ but generally γ_T varies around \sqrt{n} (see table 5.3).

To answer the second question, we will introduce the following definition:

Definition 5.3. USP coefficient

Given a timetable T and a pair of cities x, y , the USP coefficient $\tau_T(x, y) = |usps_T(x, y)|$, where $usps_T(x, y)$ is the set of USPs between x and y . By τ_T we will denote the average USP coefficient in timetable T .

From the table 5.2 we can see, that there are not many USPs on average, meaning that τ is usually some small number. Also, we see that it slightly increases with increasing time range (plot 5.3), but not with increasing n , the size of the timetable (plot 5.4). Thus we can consider τ to be bound by a small constant when it comes to daily timetables.

From the answers to our two questions we see that the **size of the preprocessed oracle** is $\mathcal{O}(\tau n^2 \gamma)$.

Name	τ	$\max \tau(x, y)$
<i>air01-200d</i>	5.8	30
<i>cpru-200d</i>	7.0	64
<i>cpza-200d</i>	5.1	42
<i>montr-200d</i>	4.3	30
<i>sncf-200d</i>	4.3	24
<i>sncf-inter-200d</i>	0.6	19
<i>sncf-ter-200d</i>	6.1	33
<i>zsr-200d</i>	2.5	19

Table 5.2: Average and maximal USP coefficients.

Name	avg USP size
<i>air01-200d</i>	3.0
<i>cpru-200d</i>	13.8
<i>cpza-200d</i>	11.1
<i>montr-200d</i>	20.3
<i>sncf-200d</i>	10.5
<i>sncf-inter-200d</i>	7.9
<i>sncf-ter-200d</i>	10.8
<i>zsr-200d</i>	13.7

Table 5.3: Average USP sizes vary around $\sqrt{n} \approx 14$. Note extremely low value for airline timetable - this is due to the fact that UGs of airline timetables have small-world characteristics [?].




Figure 5.3: Changing of τ with increased time range in *air01* dataset. 1 day = about 800 in height.




Figure 5.4: Changing of τ with increased number of stations in *sncf* dataset.

The query time also depends on the USP coefficient of a given pair of cities x, y , as we have to try out all USPs in $usps(x, y)$. The expansion of a USP by *Expand* function takes time linear in the size of the USP¹⁰, leading to **query time** $\mathcal{O}(\tau\gamma)$ on average. Note, that this is pretty much optimal, as τ is basically constant and we need to output the connection itself, which takes linear time in its size.

Finally, the **stretch** of *USP-OR* is **1**, as it returns exact answers.

<i>USP-OR</i>	<i>prep</i>	<i>size</i>	<i>qtime</i>	<i>stretch</i>
guaranteed	$\mathcal{O}(hn^2(\log n + \delta))$	$\mathcal{O}(\tau n^2 \gamma)$	avg. $\mathcal{O}(\tau\gamma)$	1
τ const., $\gamma \leq \sqrt{n}$, $\delta \leq \log n$	$\mathcal{O}(hn^2 \log n)$	$\mathcal{O}(n^{2.5})$	avg. $\mathcal{O}(\sqrt{n})$	1

Table 5.4: The summary of the *USP-OR* algorithm parameters. The second row corresponds e.g. to the *sncf* dataset.

¹⁰In time-dependent graphs, this requires a constant-time retrieval of the correct interpolation point of the cost function (the piece-wise linear function that tells us the traversal time of an arc at a given time) for some time t . More specifically, we need to obtain an interpolation point $\operatorname{argmin}_{(t', l)} \{t' \mid t' > t\}$. If we assume uniform distribution of departures throughout the time range of the timetable, this can be implemented in constant time. Otherwise, binary search lookup is possible in time $\mathcal{O}(\log h)$

5.1.2 Performance of *USP-OR*




Figure 5.5: *USP-OR* algorithm compared to TD Dijkstra on the *cpvu* dataset.




Figure 5.6: *USP-OR* algorithm compared to TD Dijkstra on the *cpza* dataset.




Figure 5.7: *USP-OR* algorithm compared to TD Dijkstra on the *montr* dataset.




Figure 5.8: *USP-OR* algorithm compared to TD Dijkstra on the *sncf* dataset.

5.2 *USP-OR-A*

With *USP-OR* the main disadvantage is its space consumption. We may decrease this space complexity by pre-computing USPs only among *some* cities. The nodes that we select for this purpose will be called **access nodes** (AN for short), as for each city they would be the crucial nodes we need to pass in order to access most of the cities of T . It would be suitable for this access node set to have several desirable properties. In order to formulate them, we need to define a few terms first.

Definition 5.4. *Front neighbourhood*

Given a timetable T and access node set \mathcal{A} , a *front neighbourhood* of city x are all cities (including x) that are reachable from x not via \mathcal{A} . Formally $\text{neigh}_{\mathcal{A}}(x) = \{y \mid \exists \text{ path } p = (p_1, p_2, \dots, p_k) \text{ from } x \text{ to } y \text{ in } ug_T : p_i \neq a \forall a \in \mathcal{A}, i \in \{2, \dots, k-1\}\}$ ¹¹

¹¹We leave out subscript identifying the timetable T . In situation with clear context, we may also leave out the \mathcal{A} subscript




Figure 5.9: *USP-OR* algorithm compared to TD Dijkstra on the *sncf-inter* dataset.




Figure 5.10: *USP-OR* algorithm compared to TD Dijkstra on the *sncf-ter* dataset.

plot:uspor-sncf-ter-size

Figure 5.11: For

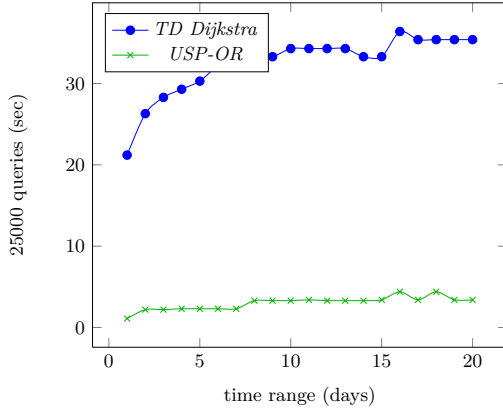


Figure 5.12: *USP-OR* algorithm compared to TD Dijkstra on the *zsr* dataset. Here we measured how increased time range influence the query time of *USP-OR*. You may see that for both algorithms the query time almost stops increasing at some point - this is because (informally) adding time range no longer brings along new optimal connections (or underlying shortest paths in case of *USP-OR*).

Name	n	spd
<i>cpru*</i>	700	14.5
<i>cpza*</i>	700	14.3
<i>montr</i>	217	8.8
<i>sncf*</i>	1000	64.8
<i>sncf-inter</i>	366	27.0
<i>sncf-ter*</i>	1000	78.3
<i>zsr</i>	233	19.3

Table 5.5: Speed-up of the *USP-OR* algorithm for the whole timetables (for those marked with asterisk we took only a subset of n stations, as we were limited by the space).

We define analogically **back neighbourhood** (denoted $bneigh_{\mathcal{A}}(x)$), as nodes that could be reached in reversed UG ($\overleftarrow{ug_T}$). Note that the access nodes that are on the boundary of x 's neighbourhoods are also part of these neighbourhoods. These access nodes form some sort of separator between the x 's neighbourhood and the rest of the graph and we will call them **local access nodes (LAN)** ($lan_{\mathcal{A}}(x) = \mathcal{A} \cap neigh_{\mathcal{A}}(x)$), or analogically **back local access nodes (blan $_{\mathcal{A}}(x)$)**.

Now we may formulate the three desired properties of the access node set \mathcal{A} . Given a timetable T and small constants r_1 , r_2 and r_3 , we would like to find access node set \mathcal{A} such that:

1. The access node set is sufficiently small

$$|\mathcal{A}| \leq r_1 \cdot \sqrt{n} \quad (5.1)$$

2. The average square of neighbourhood ¹² size for cities not in \mathcal{A} is at most $r_2 \cdot n$

$$\frac{\sum_{x \in ct_T \setminus \mathcal{A}} |neigh_{\mathcal{A}}(x)|^2}{|ct_T \setminus \mathcal{A}|} \leq r_2 \cdot n \quad (5.2)$$

3. The number of local access nodes for each node is bounded by r_3 :

$$|lan_{\mathcal{A}}(x)| \leq r_3, \forall x \in ct_T \quad (5.3)$$

An access node set \mathcal{A} with the above mentioned properties will be called **(r_1, r_2, r_3) access node set** (AN set). We will now explain how the *USP-OR-A* (USP-OR with access nodes) algorithm works and return to its analysis later.

During preprocessing, we need to find a good AN set and compute the USPs between every pair of access nodes. For every city $x \notin \mathcal{A}$, we also store its $neigh_{\mathcal{A}}(x)$, $bneigh_{\mathcal{A}}(x)$, $lan_{\mathcal{A}}(x)$ and $blan_{\mathcal{A}}(x)$. On a query from x to y at time t , we will first make a local search in the neighbourhood of x up to x 's local access nodes. Subsequently, we want to find out the earliest arrival times to each of y 's *back* local access nodes. To do this, we take advantage of the pre-computed USPs between access nodes - try out all the pairs $u \in lan(x)$ and $v \in blan(y)$ and expand the stored USPs. Finally, we make a local search from each of y 's back LANs to y , but we run the search *restricted* to y 's back neighbourhood. For more details, see algorithms 3 and 4 and picture 5.13, where we have split the algorithms to 3 distinct phases.

Algorithm 3 *USP-OR-A* preprocessing

Input

- timetable T

Algorithm

find a good AN set \mathcal{A}

$\forall x, y \in \mathcal{A}$ compute $usps(x, y)$

$\forall x \in ct_T \setminus \mathcal{A}$ compute $neigh_{\mathcal{A}}(x)$, $bneigh_{\mathcal{A}}(x)$, $lan_{\mathcal{A}}(x)$ and $blan_{\mathcal{A}}(x)$

Output

- output everything we have computed
-

¹²We required the same for back neighbourhoods

Algorithm 4 *USP-OR-A* query

Input

- timetable T
- OC query (x, t, y)

Algorithm

let $lan(x) = x$ if $x \in \mathcal{A}$

let $blan(y) = y$ if $y \in \mathcal{A}$

Local front search

perform TD Dijkstra from x at time t up to $lan(x)$

if $y \in neigh(x)$ **then**

 let c_{loc}^* be the connection to y obtained by TD Dijkstra *# the optimal connection may still go via ANs (though it is unlikely)*

end if

$\forall u \in lan(x)$ let $ea(u)$ the arrival time and $oc(u)$ the conn. to u obtained by TD Dijkstra

Inter-AN search

for all $v \in blan(y)$ **do**

$oc(v) = null$

for all $u \in lan(x)$ **do**

for all $p \in usps(u, v)$ **do**

$c = Expand(T, p, ea(u))$

$oc(v) = \text{better out of } oc(v) \text{ and } c$

end for

end for

end for

$\forall v \in blan(y)$ let $ea(v) = end(oc(v))$

Local back search

for all $v \in blan(y)$ **do**

 perform TD Dijkstra from v at time $ea(v)$ to y restricted to $bneigh(y)$

 let $fin(v)$ be the connection returned by Dijkstra's algorithm

end for

$v^* = \operatorname{argmin}_{v \in blan(y)} \{end(fin(v))\}$

$u^* = from(oc(v^*))$

let $c^* = oc(u^*).oc(v^*).fin(v^*)$ *# the dot (.) symbol is concatenation of connections*

output better out of c_{loc}^* and c^*

Output

- optimal connection $c_{(x,t,y)}^*$
-

5.2.1 Analysis of *USP-OR-A*

Let us now analyse the properties of this oracle-based method. Clearly, much depends on the way we look for the access node set. We will address this issue in next subsections but for now, we will assume we can find (r_1, r_2, r_3) AN set \mathcal{A} in time $f(n)$. Then, in the preprocessing, we have to find USPs among the access nodes, which requires running Dijkstra's algorithm from each event in a city from \mathcal{A} . There is $\mathcal{O}(r_1 h \sqrt{n})$ such events which leads to the time complexity $\mathcal{O}(r_1 h n^{1.5} (\log n + \delta))$. We also have to find local access nodes and neighbourhoods for each city, which can be accomplished with e.g. depth first search exploring the neighbourhood. This search algorithm (run from non-access city) has complexity linear in the number of arcs and so we could bound the total complexity as:

$$\sum_{x \in ct_T \setminus \mathcal{A}} |E(neigh_{\mathcal{A}}(x))| \leq \sum_{x \in ct_T \setminus \mathcal{A}} |neigh_{\mathcal{A}}(x)|^2 \leq r_2 n^2$$

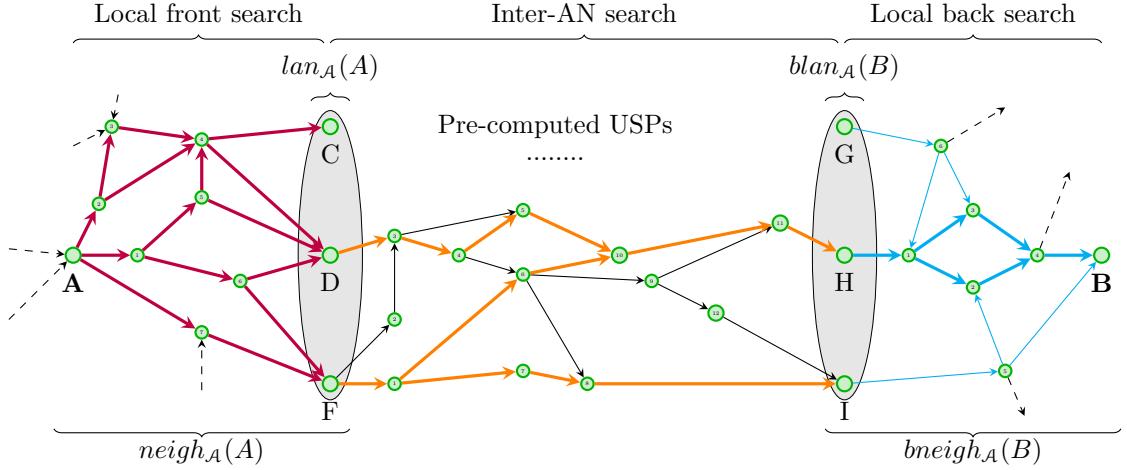


Figure 5.13: Principle of *USP-OR-A* algorithm. The arcs in **bold** mark areas that will be explored: all nodes in $neigh_A(x)$, USPs between LANs of x and back LANs of y and the back neighbourhood of y (possibly only part of it will be explored, since the local back search goes against the direction in which the back neighbourhood was created).

where $E(V)$ is the set of arcs among vertices of V . However this is very loose upper bound, as our UGs are actually very sparse. Therefore we can improve it. We know from the equation 5.2 that the average square of neighbourhood size is $\leq r_2 \cdot n$. As a consequence of the Cauchy-Schwarz Inequality [?] the following holds for positive real numbers x_i :

$$\sqrt{\frac{x_1^2 + x_2^2 + \dots + x_n^2}{n}} \geq \frac{x_1 + x_2 + \dots + x_n}{n}$$

Applying this to our neighbourhood sizes, we get that the average size of the neighbourhood is at most $\sqrt{r_2 n}$. We now split the vertices of $ct_T \setminus \mathcal{A}$ to two categories: those with neighbourhoods of size $\leq \sqrt[4]{n}$ will be part of the set S_{\leq} and those with neighbourhoods of size bigger then $\sqrt[4]{n}$ will be in $S_{>}$. A neighbourhood in the first category cannot possibly contain more than \sqrt{n} arcs while those in the second category can have at most $\delta_T |neigh_A(x)|$ arcs, depending on the timetable's density.

$$\begin{aligned} \sum_{x \in ct_T \setminus \mathcal{A}} |E(neigh_A(x))| &\leq \\ \sum_{x \in S_{\leq}} \overbrace{|E(neigh_A(x))|}^{\leq \sqrt{n}} + \sum_{x \in S_{>}} \overbrace{|E(neigh_A(x))|}^{\leq \delta |neigh_A(x)|} &\leq \\ n\sqrt{n} + \delta n\sqrt{r_2 n} &\leq \\ \delta r_2 n^{1.5} \end{aligned}$$

Therefore, the total **time complexity of the preprocessing** is $\mathcal{O}(f(n) + r_1 h n^{1.5} (\log n + \delta)) + \mathcal{O}(\delta r_2 n^{1.5}) = \mathcal{O}(f(n) + (r_1 + r_2)(\delta + \log n) h n^{1.5})$.

As for the size of the preprocessed data - we need to store all the neighbourhoods, LANs and USPs between access nodes. We already know that the average size of the neighbourhood is $\leq \sqrt{r_2 n}$, thus the total size of the (front and back) neighbourhoods is $\mathcal{O}(r_2 n^{1.5})$ ¹³. This term

¹³As r_2 will be a very small constant, we may disregard the square root

bounds also the size of the pre-computed local access nodes for each node.

Finally we have the preprocessed USPs. There is at most $r_1^2 n$ pairs of access nodes and for each of them we have possibly several USPs. We will denote by $\tau_{\mathcal{A}}$ the average USP coefficient between pairs of cities from \mathcal{A} and by $\gamma_{\mathcal{A}}$ the average optimal connection size (or equivalently, USP size) between cities in \mathcal{A} . This amounts to $\mathcal{O}(r_1^2 \tau_{\mathcal{A}} \gamma_{\mathcal{A}} n)$ for storage of USPs and to a total **preprocessing size** $\mathcal{O}(r_2 n^{1.5} + r_1^2 \tau_{\mathcal{A}} \gamma_{\mathcal{A}} n)$.




Figure 5.14: Changing of $\tau_{\mathcal{A}}$ with increased number of stations in *cpza* dataset. \mathcal{A} was obtained using algorithm *locsep* we will talk about later




Figure 5.15: Changing of $\gamma_{\mathcal{A}}$ with increased number of stations in *cpza* dataset. \mathcal{A} was obtained using algorithm *locsep* we will talk about later

On a query from x at time t to y , we first perform the *local front search* (see algorithm 4). In this step we explore the neighbourhood of x with a time-dependent Dijkstra's algorithm, which takes on average time $\mathcal{O}(\sqrt{r_2 n}(\log(\sqrt{r_2 n}) + \delta))$. We then expand all the USPs between u and v such that $u \in \text{lan}(x)$ and $v \in \text{blan}(y)$, which takes on average $\mathcal{O}(r_3^2 \tau_{\mathcal{A}} \gamma_{\mathcal{A}})$. Finally, from each $v \in \text{blan}(y)$ we do a TD Dijkstra, restricted to $\text{bneigh}(y)$, leading to time complexity $\mathcal{O}(r_3 \sqrt{r_2 n}(\log(\sqrt{r_2 n}) + \delta))$.

Summing up the three terms we obtain the **query time** of $\mathcal{O}(r_2 r_3 \sqrt{n}(\log(r_2 n) + \delta) + r_3^2 \tau_{\mathcal{A}} \gamma_{\mathcal{A}})$.

Stretch of the *USP-OR-A* algorithm is **1**, as it is exact.

The resulting bounds do not look very appealing. This is because we wanted to preserve the generality - the concrete bounds will depend on what kind of properties the timetables have and what algorithm for finding the AN set is plugged in. In table 5.6, we summarize the parameters of *USP-OR-A* method and provide the bounds for a case when the properties of the timetables correspond to those we have measured in our datasets and when we have an algorithm that finds good AN set.

<i>USP-OR-A</i>	guaranteed	τ, r_1, r_2, r_3 const., $\gamma \leq \sqrt{n}$, $\delta \leq \log n$
<i>prep</i>	$\mathcal{O}(f(n) + (r_1 + r_2)(\delta + \log n)hn^{1.5})$	$\mathcal{O}(f(n) + hn^{1.5} \log n)$
<i>size</i>	$\mathcal{O}(r_2 n^{1.5} + r_1^2 \tau_{\mathcal{A}} \gamma_{\mathcal{A}} n)$	$\mathcal{O}(n^{1.5})$
<i>qtime</i>	avg. $\mathcal{O}(r_2 r_3 \sqrt{n}(\log(r_2 n) + \delta) + r_3^2 \tau_{\mathcal{A}} \gamma_{\mathcal{A}})$	avg. $\mathcal{O}(\sqrt{n} \log n)$
<i>stretch</i>	1	1

Table 5.6: The summary of the *USP-OR-A* algorithm parameters.

5.2.2 Correctness of *USP-OR-A*

Finally, we will proof the correctness of the algorithm, i.e. that it always returns the optimal connection.

Theorem 5.1. *The algorithm USP-OR-A 3 4 always returns the optimal connection.*

Proof. Let \mathcal{A} be the set of access nodes and consider a query from city x to city y at any time t . If $x \in \mathcal{A}$ and $y \in \mathcal{A}$, an optimum is returned due to lemma 5.1 (in such a case, we basically run *USP-OR* algorithm).

In the following we will assume that $y \notin \text{neigh}(x)$, which means that the optimal connection goes through some access node $u \in \text{lan}(x)$ and $v \in \text{blan}(y)$. Note that it may be that $u = v$.

What we would like to prove as a next step is that we reach the back LANs of y (or y itself if it is an access node) at the earliest arrival time. After the *local front search*, we have reached the x 's local ANs at times $ea(u) \forall u \in \text{lan}(x)$. For some local access node this value is the true earliest arrival. Let us denote the set of such local ANs as $\text{lan}^*(x)$. The crucial thing to realize is, that the optimal connection to any city out of the x 's neighbourhood will lead via some $u \in \text{lan}^*(x)$ (see picture 5.16). And because the *inter-AN search* phase finds *optimal* connections between pairs $u \in \text{lan}(x)$ and $v \in \text{blan}(y)$, it follows that for each $v \in \text{blan}(y)$ the $ea(v)$ is the earliest arrival to this city after the *inter-AN search* phase.




Figure 5.16: On the picture $\text{lan}(x) = \{G, H\}$ and $\text{blan}(y) = \{G, H, I\}$. In **thick** we have highlighted the optimal connection. The connection to H is sub-optimal after the *local front search* phase, however the optimal connection to y (and to H and I as well) leads through $\text{lan}^*(x)$ (some of x 's local access nodes to which we have an optimal connection after the *local front search*. Particularly, it goes through G).

In the *local back search* we run a TD Dijkstra search from all back LANs of y . And since this algorithm is exact and starts from each back LAN as early as possible, we get the optimal connection to y .

It remains to show that if $y \in \text{neigh}(x)$, we also get the optimal connection. In such case, we simply compare the connection that goes via access nodes and the one that was obtained solely within the neighbourhood and output the shorter one. As there are no other options, the proof is complete. \square

5.2.3 Modifications of *USP-OR-A*

Our implementation of the *USP-OR-A* algorithm uses one slight improvement, which we did not mention in its description, since it is more of an optimization technique without any theoretical guarantees on actual improvement of the running time. However, we consider it an interesting idea so we mention it at this place.

Definition 5.5. *USP tree*

Given a pair of cities x and y in a timetable T , we will call a *USP tree* the graph made out of edges of all *USPs* in $usp_T(x, y)$: $usp_T^3(x, y) = (V^3, E^3)$ where $V^3 = \{v \mid v \text{ lays on some } p \in usp_T(x, y)\}$ and $E^3 = \{(a, b) \mid (a, b) \text{ is part of some } p \in usp_T(x, y)\}$.

We could take advantage of these *USP trees* to speed up the *local front search* phase of the algorithm, where we unnecessarily explore the whole neighbourhood when we could just go along the arcs of the *USP trees*. The picture 5.17 depicts this.




Figure 5.17: Using *USP trees* (thick non-dashed arcs in $neigh_A(x)$) to decrease the explored area in *local front search*. A full neighbourhood search is done only when $y \in neigh(x)$.

The interesting thing about this is the exploitation of both - timetable and its underlying graph. While the neighbourhood of a node is something static, related only to the structure of the UG and generally time-independent, the *USP trees* reflect to some extent the properties of the timetable (e.g. which ways are frequently serviced and thus provide optimal connections). By intersecting these two things, we get the area that is *worth* to be explored and that is *small* at the same time (provided, of course, that the neighbourhoods are small).

5.3 Selection of access node set

The challenge in the *USP-OR-A* algorithm comes down to the selection of a good access node set - a (r_1, r_2, r_3) AN set with both three parameters as low as possible. However, intuitively (and experimentally verified), decreasing e.g. r_1 (the AN set size) increases r_2 (the size of the neighbourhoods). We therefore have to do some compromises.

In the following we first show the problem of choosing an optimal access node set to be NP-hard. We then present our methods for heuristic selection of access nodes and show their performance on real data.

5.3.1 Choosing the optimal access node set

A question stands - what is an optimal access node set? To keep the query time as low as possible, we need to avoid large neighbourhood sizes, because that would mean spending too much time doing local searches. A pretty good upper bound for neighbourhood sizes seems to be \sqrt{n} (i.e. $r_1 = 1$) - the idea is that in such case the local searches cannot possibly last longer than $\mathcal{O}(n)$ while the *inter-AN search* is linear in the size of the connection and can also be at most $\mathcal{O}(n)$. In practice, both of these steps will be faster because the neighbourhoods are sparse and because the connections are on average much shorter than n . However, it gives an idea of why \sqrt{n} should be considered for a target neighbourhood size.

Therefore, the question stands: What is the smallest set of ANs, such that the neighbourhood sizes are all under \sqrt{n} ? More formally, for a timetable T , the task is to minimize $|\mathcal{A}|$ where $\mathcal{A} \subseteq ct_T$ and $\forall x \in ct_T \setminus \mathcal{A} : |neigh_{\mathcal{A}}(x)| \leq \sqrt{n}$. We will call this the **problem of the optimal access node set** and in what follows we will show that it is NP-complete.

Theorem 5.2. *The problem of the optimal access node set is NP-complete*

Proof. We will make a reduction of the *min-set cover* problem (a NP-complete problem) to the problem of optimal AN set.

Consider an instance of the min-set cover problem:

- A universe $U = \{1, 2, \dots, m\}$
- k subsets of U : $S_i \subseteq U$ $i = \{1, 2, \dots, k\}$ whose union is U : $\bigcup_{1 \leq i \leq k} S_i = U$

Denote $\mathcal{S} = \{S_i \mid 1 \leq i \leq k\}$. The task is to choose the smallest subset \mathcal{S}^* of \mathcal{S} that still covers the universe ($\bigcup_{S_i \in \mathcal{S}^*} S_i = U$). We will now do a simple conversion (in polynomial time) of the instance of min-set cover to the instance of the optimal AN set problem (which is represented by the underlying graph of T).

For each $j \in U$, we will make a complete graph of β_j vertices (the value of β_j will be discussed later) named m_j and for each set S_i we will make a vertex s_i and vertex s'_i . We will connect all vertices of m_j to $s_i \iff j \in S_i$. Finally, for we will connect s_i to s'_i for $1 \leq i \leq k$.

Example. Let $m = 10$ (thus $U = \{1, 2, \dots, 10\}$) and $k = 13$:

- $S_1 = \{1, 3, 10\}$
- $S_2 = \{1, 2\}$
- ...
- $S_{13} = \{2, 3, 10\}$

For this instance of min set-cover, we construct the graph depicted on picture 5.18.

Now we would like to clarify the sizes of m_i . Define α_i to be the number of sets S_j that contain i : $\alpha_i = |\{S_j \in \mathcal{S} \mid i \in S_j\}|$ and assume the constructed graph has n vertices. We want the β_i to satisfy $\beta_i \geq 2$ and $\beta_i + 2\alpha_i - 1 \leq \sqrt{n}$ but $\beta_i + 2\alpha_i > \sqrt{n}$. The last two inequalities would mean that if at least one s_j connected to m_i is chosen as an access node, the neighbourhood for nodes in m_i will be still $\leq \sqrt{n}$, but if none of them is chosen, the neighbourhood will be just over \sqrt{n} . For now we will assume that we have constructed the graph in such a way that all β_i satisfy the mentioned inequalities. We will return to construction of the graph at the end.

Now consider an optimal AN set which contains a vertex from within some m_i . If this is the case, **either** some s_j to which m_i is connected is selected as AN, **or** all vertices from m_i are




Figure 5.18: The principle of the reduction. In m_i , there are actually complete graphs of β_i vertices (as shown for m_1). **Thick** arcs represent arcs from all the vertices of respective m_i . The s_i vertices are connected to their s'_i versions. If e.g. s_1 is selected as an access node, s'_1 is no longer part of any neighbourhood.

access nodes **or** the neighbourhood is too large. Keep in mind that the local access nodes are also part of neighbourhoods, so unless we select for AN some of the s_j that m_i is connected to, the neighbourhood of any non-access node in m_i will be too large. As there are at least two nodes in every m_i , it is more efficient to select some s_j rather than select all nodes in m_i . Thus when it comes to selecting ANs *it is worth to consider only vertices s_j* .

From this point on, it is easy to see that it is optimal to select those s_j that correspond to the optimal solution of min-set cover. The reason is that each of the m_i will be connected to at least one access node s_j and will thus have neighbourhood size $\leq \sqrt{n}$, while the number of selected access nodes will be optimal.

It remains to show how to choose values β_i . Due to the condition $\beta_i \leq \sqrt{n} - 2\alpha_i + 1$ we need to have sufficiently big n to fulfil $\beta_i \geq 1$. We will accomplish this by adding dummy isolated vertices to the graph. Define function $nextSquare(x)$ to output the smallest $y^2 > x$ where y is a natural number. We then compute $w = (\max\{2\alpha_i\} + 2)^2$ and select the starting value of n to be $n' = nextSquare(\max\{w - 1, \sqrt{2k + m}\})$. We create the s_j and s'_j vertices and complete graphs m_i containing so far only one vertex each. We connect everything according to the rules stated earlier in this proof and we create dummy vertices up to the capacity defined by n . Now we repeat the following:

- We compute \sqrt{n} which is a natural number
- For i from 1 to m we add vertices to m_i till it does not contain $\sqrt{n} - 2\alpha_i + 1$ vertices. For each added vertex we delete one dummy vertex.
- If we run out of dummy vertices, $n = nextSquare(n)$
- Break out of the loop if $|m_i| = \sqrt{n} - 2\alpha_i + 1 \forall i$

With each iteration of this little algorithm we will be forced to add one more vertex to all m_i (since \sqrt{n} increased by one), a so called *inefficient increase*. At the beginning, we need to make at most $m\sqrt{n'}$ efficient increases to meet the breaking condition. And since m is constant and the capacity of new dummy vertices increases linearly, after t steps we create $\mathcal{O}(t^2)$ dummy vertices that may be used for efficient increases. Therefore, the algorithm will stop after $\mathcal{O}(\sqrt{mn'})$ steps. \square

5.3.2 Choosing ANs based on node properties

In the previous sub-subsection, we have shown the problem of choosing the optimal AN set to be NP-hard. In this sub-subsection we perform a simple experiment of choosing for the access nodes the cities that seem to be the most important. More specifically, in the optimistic underlying graph (see section 2) ug_T^{opt} we were looking for cities with:

1. High **degree**. We consider the sum of in-degree and out-degree¹⁴ of the respective node x :
 $deg(x) = deg_{in}(x) + deg_{out}(x)$.
2. High **betweenness centrality** (BC). Betweenness centrality for a node v is defined as

$$g(v) = \sum_{s \neq v \neq t} \frac{\sigma_{st}(v)}{\sigma_{st}}$$

where $\sigma_{st}(v)$ is the number of shortest paths from s to t passing through v and σ_{st} is the total number of shortest paths from s to t . [?]. We then scale the values to the range $< 0, 1 >$ to obtain for each city x its scaled betweenness centrality $bc(x)$.

We will denote by $\mathcal{A}_{deg}(k)$ the set of k cities with highest $deg(x)$ value and similarly $\mathcal{A}_{bc}(k)$ will be the set of k cities with the highest $bc(x)$ value. We were interested in the smallest k such that:

1. $\mathcal{A}_{deg}(k)$ is (r_1, r_2, r_3) AN set with $r_2 \leq 1$ (the average square of neighbourhoods is $\leq \sqrt{n}$). Denote such k as k_{deg}^{avg} .
2. $\forall x \in c_T: |neigh_{\mathcal{A}_{deg}(k)}(x)| \leq \frac{3\sqrt{n}}{2}$ and $|bneigh_{\mathcal{A}_{deg}(k)}(x)| \leq \frac{3\sqrt{n}}{2}$ (all neighbourhoods are large at most $\frac{3\sqrt{n}}{2}$). Denote such k as k_{deg}^{max} .

We define similarly k_{bc}^{avg} and k_{bc}^{max} . The resulting values for datasets *sncf* and *cpru* could be seen on plots 5.19.

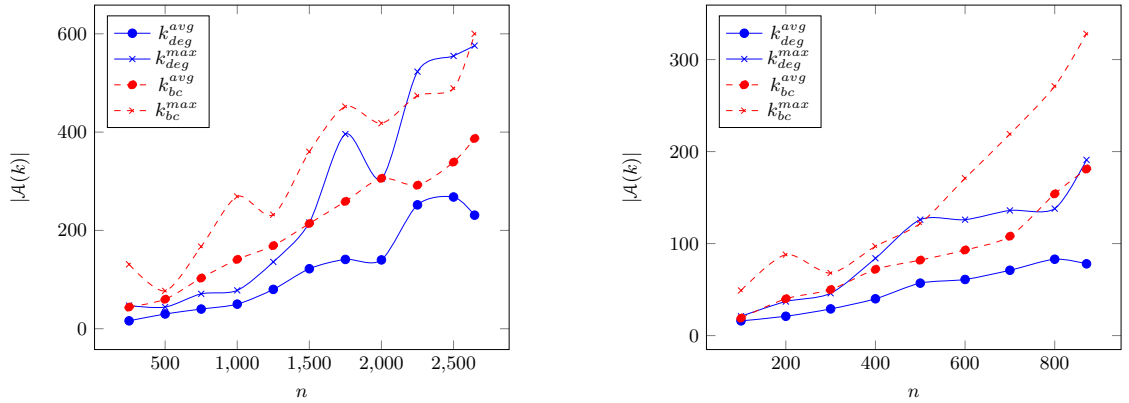


Figure 5.19: Necessary sizes of access node sets based on high-degree/high-BC cities. Datasets *sncf* (left) and *cpru* (right). Notice the occasional “roller coaster” bumps (especially on the left plot) - an explanation of this phenomena is that in the immediately smaller sub-timetable we have erased just the high-degree node that proved to be a good access node, and which now must be substituted by many other access nodes.

¹⁴In-degree is the number of arcs going towards to node and out-degree the number of outgoing arcs.

5.3.3 Choosing ANs heuristically

Clearly, selecting the cities for access nodes solely by high degree or BC value is not the best way. Probably the few nodes with highest degrees and BC will indeed be part of the AN set, as they are intuitively some sort of central hubs without which the network would not work. However, after we select the these most important nodes to the AN set, we need some better measure of node's importance, or suitability to be an access node. In the following we present a simple heuristic approach run on underlying graph ug_T of given timetable T that evaluates its vertices based on how good local separators they are.

The algorithm will work in iterations, each of them resulting in a selection ¹⁵ of a city with highest score to the access node set \mathcal{A} . Similarly to the previous approach that used degree/BC, we consider two stopping criterion:

1. \mathcal{A} is (r_1, r_2, r_3) AN set with $r_2 \leq 1$ (the average square of neighbourhoods is $\leq \sqrt{n}$). Denote $k_{locsep}^{avg} = |\mathcal{A}|$ when this stopping criterion is met.
2. $\forall x \in ct_T: |neigh_{\mathcal{A}}(x)| \leq \frac{3\sqrt{n}}{2}$ and $|bneigh_{\mathcal{A}}(x)| \leq \frac{3\sqrt{n}}{2}$ (all neighbourhoods are large at most $\frac{3\sqrt{n}}{2}$). Denote $k_{locsep}^{max} = |\mathcal{A}|$ when this stopping criterion is met.

The algorithm using the first stopping criterion will be referred to as *Locsep* and the one using the second criterion as *Locsep Max*. The algorithms differ only in the stopping criterion, therefore what follows is common to both of them.

In each iteration, we compute for each city x its score, which is done in the following way: we explore an area A_x of \sqrt{n} nearest cities around x . We do this in an underlying graph with no orientation and no weights. Next we get the front and back neighbourhoods of x within A_x ($fn(x) = neigh(x) \cap A_x$, $bn(x) = bneigh(x) \cap A_x$).

For a set of access nodes \mathcal{A} , let us call a path p in ug_T **access-free** if it does not contain a node from \mathcal{A} . Now as long as x is not in \mathcal{A} , there is a guarantee that for every pair $u \in bn(x)$ and $v \in fn(x)$ there is an access-free path from u to v within A_x . Our interest is how this will change after the selection of x . Let us call $af(x)$ the number of pairs $u \in bn(x)$ and $v \in fn(x)$ such that there is an access-free path within A_x after the selection of x to \mathcal{A} . We define the **potential** of x as:

$$p_x = \frac{af(x)}{|bn(x)||fn(x)|}$$

Our algorithm will compute a lower bound on p_x .

5.4 Performance of *USP-OR-A* and comparisons

In this subsection we give the results of *USP-OR-A* performance on our datasets.

5.4.1 *USP-OR-A* with *Locsep*

On picture

5.4.2 *USP-OR-A* with *Locsep Max*

¹⁵Actually, in our implementation, we allow an occasional de-selection of an already selected node with the *lowest* score, to avoid having in the resulting set cities that had high score when selected but were not very useful access nodes at the end.




Figure 5.20: The principle of computing potentials in Locsep algorithm. We explored an area of \sqrt{n} nearest cities (in terms of hops) around x . Little **squares** are nodes from $neigh(x)$ and **diamonds** are part of $bneigh(x)$. The highlighted area represents the back neighbourhood for node b_2 . From its nodes we run a forward search (the **thick** arcs). Nodes from the $neigh(x)$ that were not explored in this search can only be reached via x itself.




Figure 5.21: *USP-OR-A* algorithm with *Locsep* compared to TD Dijkstra on the *cpru* dataset.




Figure 5.22: *USP-OR-A* algorithm with *Locsep* compared to TD Dijkstra on the *cpza* dataset.

6 Neural network approach




Figure 5.23: *USP-OR-A* algorithm with *Loc-sep* compared to TD Dijkstra on the *montr* dataset.




Figure 5.24: *USP-OR-A* algorithm with *Loc-sep* compared to TD Dijkstra on the *sncf* dataset.




Figure 5.25: *USP-OR-A* algorithm with *Loc-sep* compared to TD Dijkstra on the *sncf-inter* dataset.




Figure 5.26: *USP-OR-A* algorithm with *Loc-sep* compared to TD Dijkstra on the *sncf-ter* dataset.

7 Application TTBlazer




Figure 5.27: *USP-OR-A* algorithm with *Locsep* compared to TD Dijkstra on the *zsr* dataset. Here we measured how increased time range influence the query time of *USP-OR-A*.




Figure 5.28: *USP-OR-A* algorithm with *Locsep Max* compared to TD Dijkstra on the *cpru* dataset.

Name	n	spd	r_1
<i>cpru</i>	871	1.7	1.5
<i>cpza</i>	1108	1.7	1.7
<i>montr</i>	217	1.5	1.3
<i>sncf</i>	2646	5.4	1.9
<i>sncf-inter</i>	366	3.6	1.5
<i>sncf-ter</i>	2637	6.3	1.6
<i>zsr</i> (daily)	233	2.14	1.1

Table 5.7: Speed-up of the *USP-OR-A* algorithm with *Locsep* for the whole timetables. In the last column is the r_1 parameter of \mathcal{A} ($|\mathcal{A}| = r_1\sqrt{n}$).




Figure 5.29: *USP-OR-A* algorithm with *Locsep Max* compared to TD Dijkstra on the *cpza* dataset.

8 Conclusion




Figure 5.30: *USP-OR-A* algorithm with *Loc-sep Max* compared to TD Dijkstra on the *montr* dataset.




Figure 5.31: *USP-OR-A* algorithm with *Loc-sep Max* compared to TD Dijkstra on the *sncf* dataset.




Figure 5.32: *USP-OR-A* algorithm with *Loc-sep Max* compared to TD Dijkstra on the *sncf-inter* dataset.




Figure 5.33: *USP-OR-A* algorithm with *Loc-sep Max* compared to TD Dijkstra on the *sncf-ter* dataset.

Appendices

A File formats

Timetable is simply a set of elementary connections, thus the format is:

- number of el. connections
- the list of all el. connections (one per line, format “*FROM TO DEP-DAY DEP-TIME ARR-DAY ARR-TIME*”)

```

1 7 //number of elementary connections
2 A B 0 10:00 0 10:45 //el. connection
3 A B 0 11:00 0 11:45
4 A B 0 12:00 0 12:45
5 A C 0 09:30 0 10:00
6 A C 0 10:15 0 10:45
7 C D 0 11:00 0 11:30
8 C D 0 13:00 0 13:30

```





Figure 5.34: *USP-OR-A* algorithm with *Locsep Max* compared to TD Dijkstra on the *zsr* dataset. Here we measured how increased time range influence the query time of *USP-OR-A*.

Name	n	spd	r_1
<i>cpru</i>	871	2.0	2.1
<i>cpza</i>	1108	1.7	2.1
<i>montr</i>	217	2.2	3.3
<i>sncf</i>	2646	6.2	2.2
<i>sncf-inter</i>	366	4.7	1.4
<i>sncf-ter</i>	2637	7.6	1.8
<i>zsr</i> (daily)	233	2.4	1.5

Table 5.8: Speed-up of the *USP-OR-A* algorithm with *Locsep Max* for the whole timetables. In the last column is the r_1 parameter of \mathcal{A} ($|\mathcal{A}| = r_1\sqrt{n}$).

Listing 1: TT file format.

Underlying graph is basically an oriented graph, with some optional parameters. The format is the following:

- number of cities
- number of arcs
- the list of all cities (one per line)
 - optional coordinates (otherwise null)
- the list of all arcs (one per line, format “*FROM TO*”)
 - optional length (otherwise null)
 - optional list of lines operating on that arc (otherwise null)

```

1 4 //number of cities
2 5 //number of arcs
3 A 45 32 //name of the city, optional coordinates
4 B null
5 C 56 34
6 D null
7 A B 57 Northern //arc, optional length and list of lines
8 A C null Picadilly Victoria
9 C B 45 Circle Jubilee Picadilly
10 C D 32 null
11 D A null null

```

Listing 2: UG file format.

Time-expanded graph is simply an oriented weighted graph, with nodes being the events and arcs being the elementary connections or waiting edges:

- number of nodes (i.e. events)
- number of arcs (el. connections + waiting)
- the list of all events (in the format “*CITY DAY TIME*”)

- the list of all arcs (in the format “FROM-EVENT TO-EVENT”)

```

1 5 //number of events
2 15 //number of arcs
3 A 0 13:30 //event
4 A 0 14:00
5 B 0 13:45
6 B 0 15:00
7 C 0 14:15
8 A 0 13:30 A 0 14:00 //waiting arc
9 A 0 13:30 B 0 13:45 //el. connection arc
10 A 0 14:00 B 0 15:00
11 A 0 13:30 B 0 15:00
12 C 0 14:15 B 0 15:00
13 ...

```

Listing 3: TE file format.

Time-dependent graph is an oriented graph with a function on the arc specifying the arc’s traversal time at any moment. In timetable networks this function is piece-wise linear and it is fully represented by the list of its interpolation points. Thus the TD file format:

- number of cities
- number of arcs
- the list of all cities (one per line)
 - optional coordinates (otherwise null)
- the list of all arcs (one per line). Arc has the format “*FROM TO INT-POINTS*” where *INT-POINTS* is a list of interpolation points¹⁶, see the listing 4 for an example.

```

1 4 //number of stations
2 5 //number of arcs
3 A 0 0 //name of the city, optional coordinates
4 B 4 4
5 C null
6 D 12 0
7 A B (0 13:30 45) (0 14:00 40) //arc and the list of interpolation
   points
8 A C (1 14:15 10)
9 C B (0 15:00 20)
10 C D (2 10:00 70)
11 D A (1 17:20 35) (1 18:00 40) (1 18:50 35)
12 ...

```

Listing 4: TD file format.

¹⁶An interpolation point is described by a triple “*DAY TIME MINUTES*”, where *MINUTES* are the traversal time