


Underlying shortest paths in timetables

František Hajnovič¹

Supervisor: Rastislav Kráľovič¹

¹ Katedra informatiky, FMFI UK, Mlynská Dolina, 842 48 Bratislava



Introduction

We introduce methods to speed-up optimal connection queries in timetables based on pre-computing paths that are worth to follow. We present two methods:

- *USP-OR* - very fast, but space consuming
- *USP-OR-A* - fast, less space consuming

And we compare the methods to:

- *Time-dependent Dijkstra*¹ - slowest, but least space consuming

We define the most common terms:

- **Timetable** - a set of **elementary connections** [Müller-Hannemann et al., 2007] (see figure 1 for an example)
- **Connection** - a valid sequence of elementary connections (may include also some waiting)
- **Underlying graph** of the timetable - nodes are the cities and there is an arc between two cities if some el. connection connects them.

Place		Time	
From	To	Departure	Arrival
A	B	10:00	10:45
B	C	11:00	11:30
B	C	11:30	12:10
B	A	11:20	12:30
C	A	11:45	12:15




Figure 1: A timetable and its underlying graph

- **Underlying shortest path (USP)** - every path p in the underlying graph, such that for some optimal connection c^* : $path(c^*) = p$ (function $path$ extracts the sequence of cities visited by the connection, see figure 2).

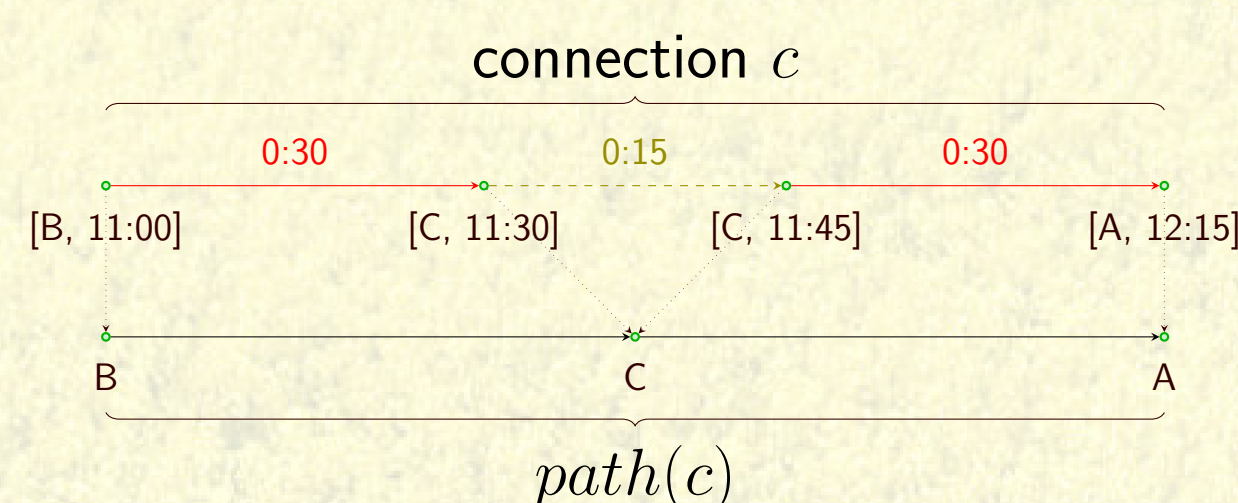


Figure 2: The $path$ function gets the underlying path from a connection.

The timetables we used for testing are real datasets and they had the following properties:

- Their underlying graphs were sparse ($m \leq n \log n$, n = number of cities)
- Average optimal connection was generally up to \sqrt{n}
- Average number of USPs between two cities was a small and constant (or slightly increasing with n)

USP-OR

Our first method, called *USP-OR* (**USP** oracle), is based on pre-computing USPs between every pair of cities. Then, upon a query from x to y at time t we consider one by one the computed USPs between x and y and perform a reverse operation to the $path$ function - $expand(p)$ where p is an USP. The $expand$ function simply follows the sequence of cities in p and from each of them it takes the first available el. connection to the next one, constructing one by one a connection from x to y .

¹Implemented in $\mathcal{O}(m + n \log n)$ using Fibonacci heap priority queues [Sommer, 2010]


<i>USP-OR</i>	<i>prep</i>	<i>size</i>	<i>qtime</i>	<i>stretch</i>
Our timetables	$\mathcal{O}(hn^2 \log n)$	$\mathcal{O}(n^{2.5})$	avg. $\mathcal{O}(\sqrt{n})$	1

Table 1: Preprocessing time/size, query time and stretch (1 = exact answers) for *USP-OR*

USP-OR-A

To decrease the space complexity, in *USP-OR-A* (**USP** oracle with access nodes) we compute USPs only among cities from a smaller set - called **access node set** (AN set \mathcal{A}). On a query from x to y at time t , we do:

- *Local front search*: a local search in the neighbourhood of x up to x 's local access nodes (LANs)
- *Inter-AN search*: expand all USPs between x 's LANs and y 's back LANs
- *Local back search*: a local search from y 's back LANs to y , restricted to y 's back neighbourhood



How to choose a good access node set?

- Solve optimally \rightarrow NP-complete (reduction of minimal set cover)
- Heuristic approach \rightarrow algorithm *Locsep* (AN set made out of nodes that are good local separators)




Figure 4: To evaluate the node's suitability to be an access node, we see what would happen to the reachability in its surroundings if we select it to the AN set.

<i>USP-OR-A + Locsep</i>	<i>prep</i>	<i>size</i>	<i>qtime</i>	<i>stretch</i>
Our timetables	$\mathcal{O}(\delta n^{2.5})$	$\mathcal{O}(n^{1.5})$	avg. $\mathcal{O}(\sqrt{n} \log n)$	1

Performance

We measured:

- **Speed-up** - how many times faster is the algorithm against the Time-dependent Dijkstra
- **Size-up** - how many times more memory is needed then to store the timetable itself

	Dataset	<i>cpsk</i> (1 day)	<i>snCF</i> (1 day)	<i>gb-coach</i> (1 week)	<i>gb-tr</i> (1 week)
<i>USP-OR</i>	n	700	1000	1000	700
	Speed-up	1	1	1	1
	Size-up	1	1	1	1
<i>USP-OR-A</i>	n	1905	2645	3915	2575
	Speed-up	1	1	1	1
	Size-up	1	1	1	1

Table 3: Speed-ups and size-ups of *USP-OR* and *USP-OR-A* with *Locsep*. Datasets from *cp.sk*, *SNCF* (French railways) and Great Britain coaches and trains.

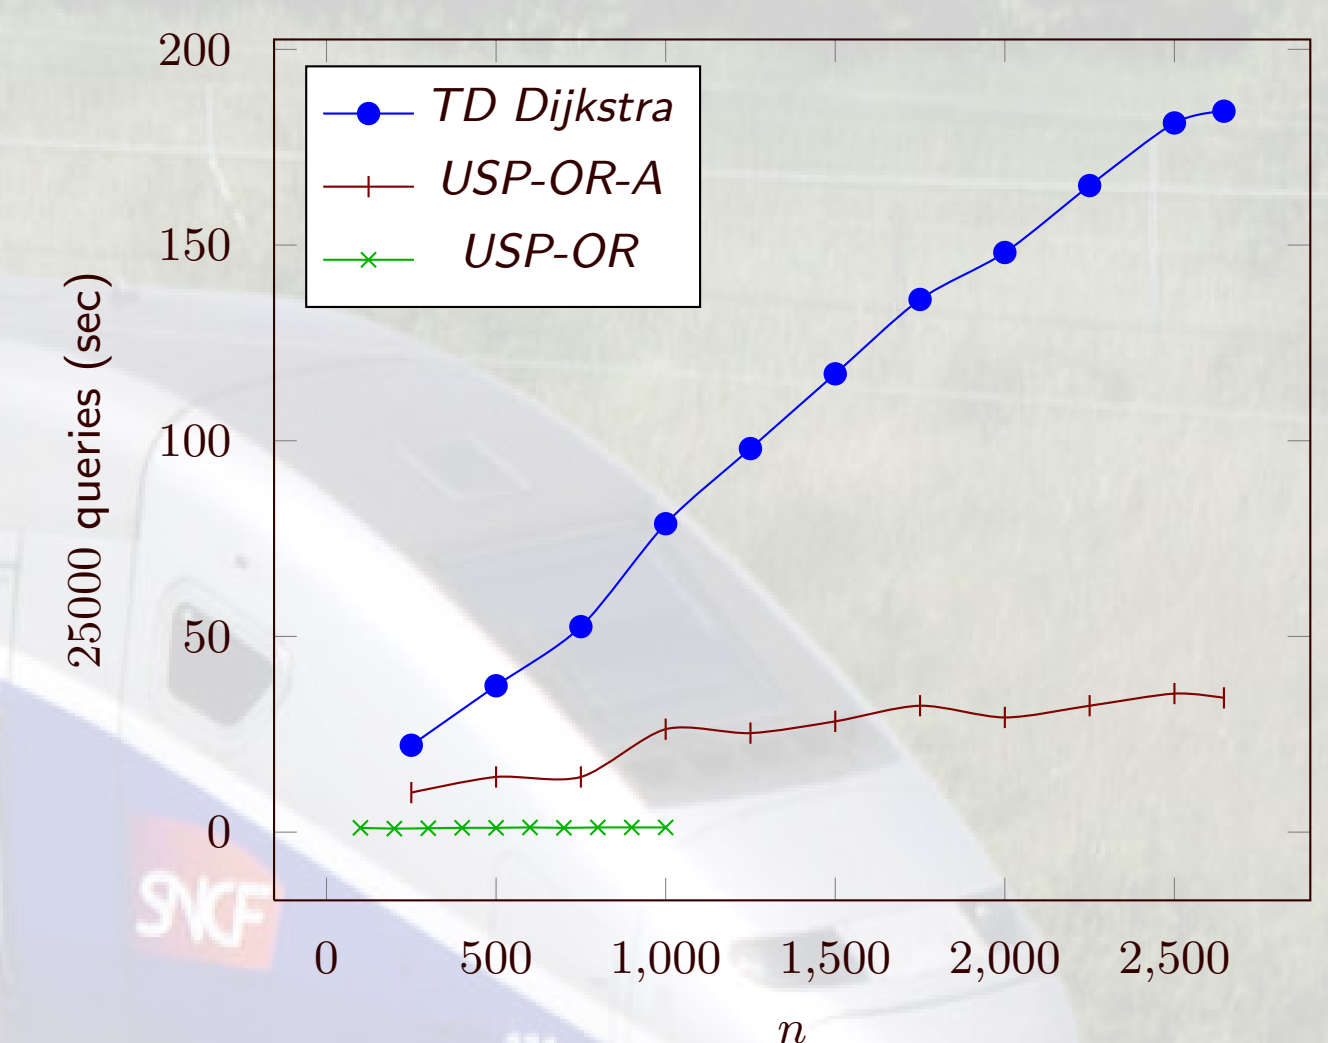


Figure 5: Query time of *USP-OR-A* with *Locsep* compared to TD Dijkstra on the *snCF* dataset. **Changing n .**

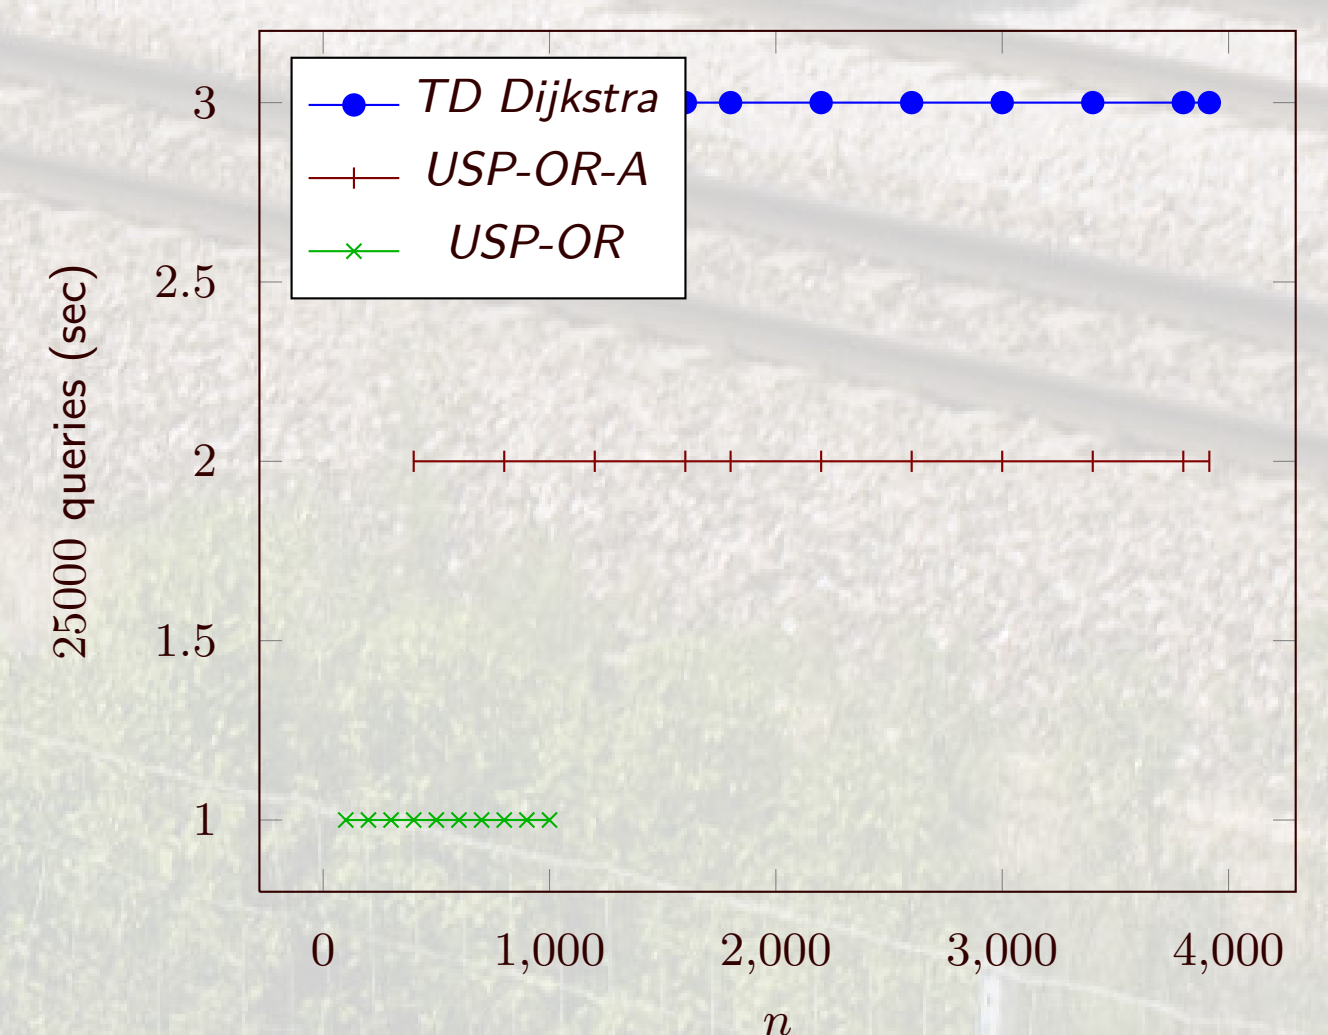


Figure 6: Query time of *USP-OR-A* with *Locsep* compared to TD Dijkstra on the *gb-coach* dataset. **Changing n .**

References

- [Müller-Hannemann et al., 2007] Müller-Hannemann, M., Schulz, F., Wagner, D., and Zaroliagis, C. (2007). *Algorithmic Methods for Railway Optimization*, volume 4359 of *Lecture Notes in Computer Science*, chapter Timetable Information: Models and Algorithms, pages 67 – 90. Springer.
- [Sommer, 2010] Sommer, C. (2010). *Approximate Shortest Path and Distance Queries in Networks*. PhD thesis, Graduate School of Information Science and Technology, The University of Tokyo.